

# modeval

November 19, 2021

## 1 Classification Model Template

```
[ ]: import joblib
df = joblib.load('finaldf.pkl')
df.head()
```

```
[ ]: id      cons_12m    cons_gas_12m  cons_last_month \
0  48ada52261e7cf58715202705a0451c9    309275          0        10025
1  24011ae4ebbe3035111d65fa7c15bc57        0        54946        0
2  d29c2c54acc38ff3c0614d0a653813dd      4660          0        0
3  764c75f661154dac3a6c254cd082ea7d      544           0        0
4  bba03439a292a1e166f80264c16191cb     1584          0        0

has_gas  nb_prod_act  num_years_antig  pow_max  active_dur  price_p1_var \
0        0            1              3   180.000     1460.0    0.151367
1        1            2              3   43.648      1096.0    0.151367
2        0            1              6   13.800      2566.0    0.151367
3        0            1              6   13.856      2192.0    0.149626
4        0            1              6   13.200      2192.0    0.149626

price_p2_var  price_p3_var  price_p1_fix  price_p2_fix  price_p3_fix  churn
0        0.0        0.0      44.266931      0.0        0.0        0
1        0.0        0.0      44.266931      0.0        0.0        1
2        0.0        0.0      44.266931      0.0        0.0        0
3        0.0        0.0      44.266931      0.0        0.0        0
4        0.0        0.0      44.266931      0.0        0.0        0

[ ]: df.drop(['id'], axis = 1, inplace = True)
df.head()
```

```
[ ]: cons_12m    cons_gas_12m  cons_last_month  has_gas  nb_prod_act \
0  309275          0        10025          0        1
1        0        54946          0          1        2
2        4660          0          0          0        1
3        544           0          0          0        1
4       1584          0          0          0        1
```

```

  num_years_antig  pow_max  active_dur  price_p1_var  price_p2_var  \
0              3  180.000     1460.0      0.151367          0.0
1              3   43.648     1096.0      0.151367          0.0
2              6   13.800     2566.0      0.151367          0.0
3              6   13.856     2192.0      0.149626          0.0
4              6   13.200     2192.0      0.149626          0.0

  price_p3_var  price_p1_fix  price_p2_fix  price_p3_fix  churn
0          0.0    44.266931        0.0        0.0       0
1          0.0    44.266931        0.0        0.0       1
2          0.0    44.266931        0.0        0.0       0
3          0.0    44.266931        0.0        0.0       0
4          0.0    44.266931        0.0        0.0       0

```

```
[ ]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import GaussianNB
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
import pandas as pd
```

```
[ ]: model_params = {
    'LogisticRegression': {
        'model': LogisticRegression(solver='liblinear', multi_class='auto'),
        'params': {
            'C': [1, 5, 10]
        }
    },
    'DecisionTree': {
        'model': DecisionTreeClassifier(),
        'params': {
            'criterion': ['gini', 'entropy']
        }
    },
    'RandomForest': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [1, 5, 10]
        }
    },
    'GaussianNB': {
        'model': GaussianNB(),
        'params': {
```

```

        }
    },
    'KNN': {
        'model': KNeighborsClassifier(),
        'params': {
            'n_neighbors': [2, 3]
        }
    }
}
X = df.drop(['churn'], axis=1)
y = df.churn

scores = []
for model_name, mp in model_params.items():
    clf = GridSearchCV(mp['model'], mp['params'],
                       cv=5, return_train_score=False)
    clf.fit(X, y)
    scores.append({
        'model': model_name,
        'best_score': (clf.best_score_).round(2),
        'best_params': (clf.best_params_),
    })
result = pd.DataFrame(scores, columns=['model', 'best_score', 'best_params'])
result

```

```
[ ]:      model  best_score          best_params
0  LogisticRegression     0.90           {'C': 1}
1      DecisionTree     0.83  {'criterion': 'gini'}
2      RandomForest     0.90  {'n_estimators': 10}
3      GaussianNB      0.23            {}
4          KNN         0.90  {'n_neighbors': 2}
```

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report,
accuracy_score, precision_score, recall_score
from matplotlib import pyplot as plt
import seaborn as sns
sc = []
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=3)

model_sel = {
    'LogisticRegression': {
        'model': LogisticRegression(solver='liblinear', multi_class='auto')
    },
    'DecisionTree': {
        'model': DecisionTreeClassifier(criterion='gini',
                                         max_depth=5,
                                         min_samples_leaf=2)
    },
    'RandomForest': {
        'model': RandomForestClassifier(n_estimators=10,
                                         max_depth=5,
                                         min_samples_leaf=2)
    },
    'GaussianNB': {
        'model': GaussianNB()
    },
    'KNN': {
        'model': KNeighborsClassifier(n_neighbors=3)
    }
}
```

```

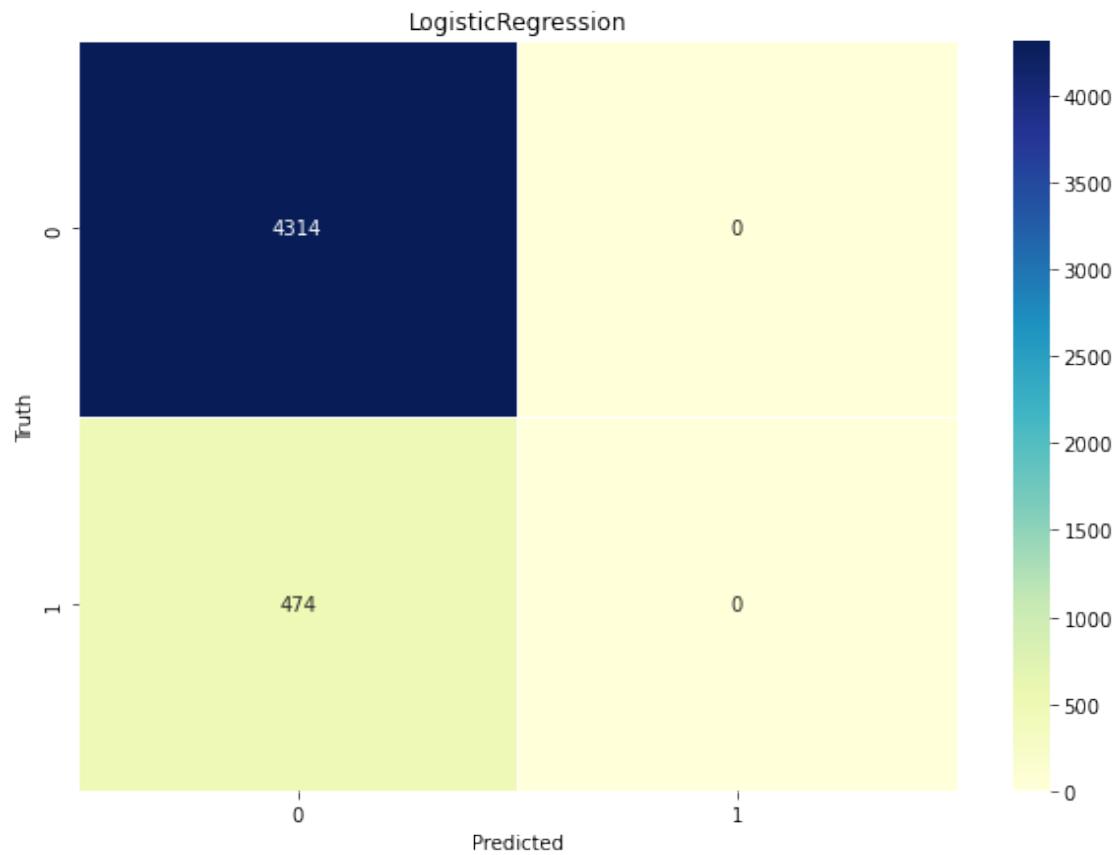
'DecisionTree': {
    'model': DecisionTreeClassifier()
},
'RandomForest': {
    'model': RandomForestClassifier()
},
'GaussianNB': {
    'model': GaussianNB()
},
'KNN': {
    'model': KNeighborsClassifier()
}
}
scor = []
for model_name, mp in model_sel.items():
    clf = mp['model']
    clf.fit(X_train, y_train)
    scor.append({
        'model': model_name,
        'accuracy': accuracy_score(y_test, clf.predict(X_test)).round(2),
        'precision': precision_score(y_test, clf.predict(X_test)).round(2),
        'recall': recall_score(y_test, clf.predict(X_test)).round(2),
    })
    cm = confusion_matrix(y_test, clf.predict(X_test)).round(2)
    plt.figure(figsize=(10, 7))
    sns.heatmap(cm, annot=True, fmt='d', cmap='YlGnBu', linewidths = .5)
    plt.xlabel('Predicted')
    plt.ylabel('Truth')
    plt.title(model_name)
    plt.show()
reslt = pd.DataFrame(scor, columns=['model', 'accuracy', 'precision', 'recall'])
reslt

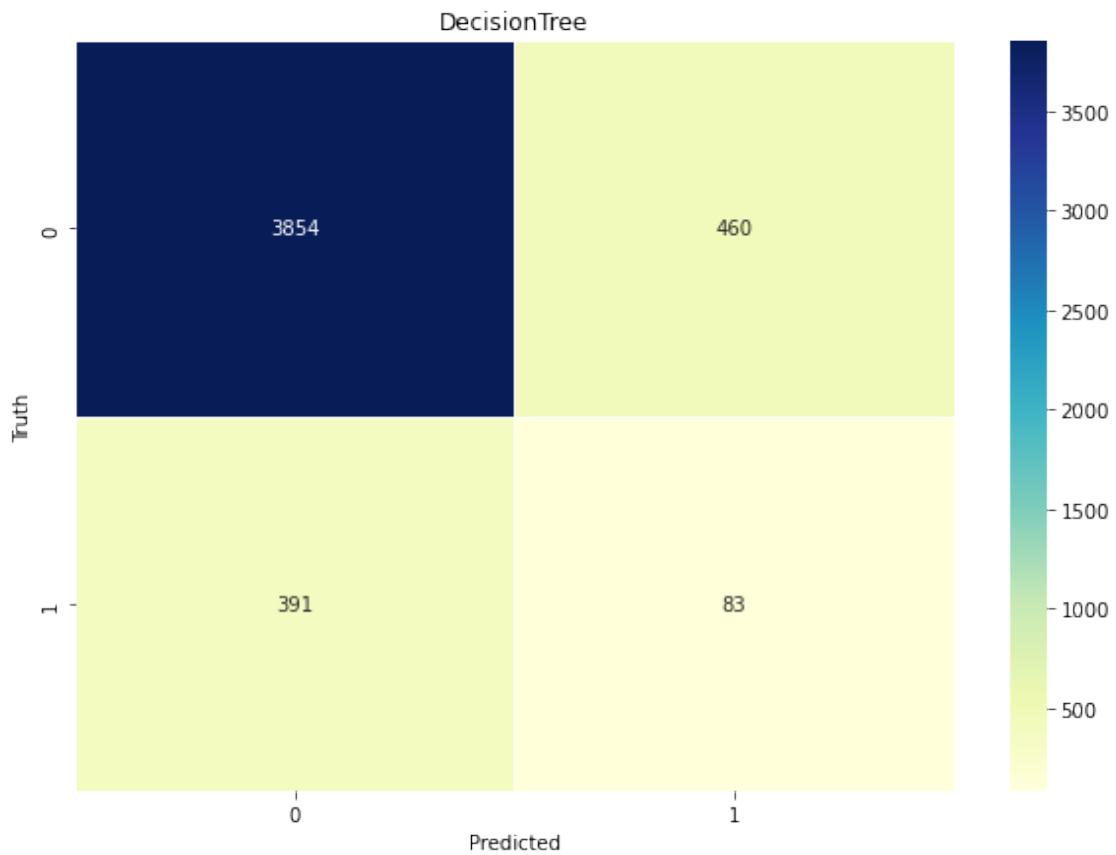
```

```

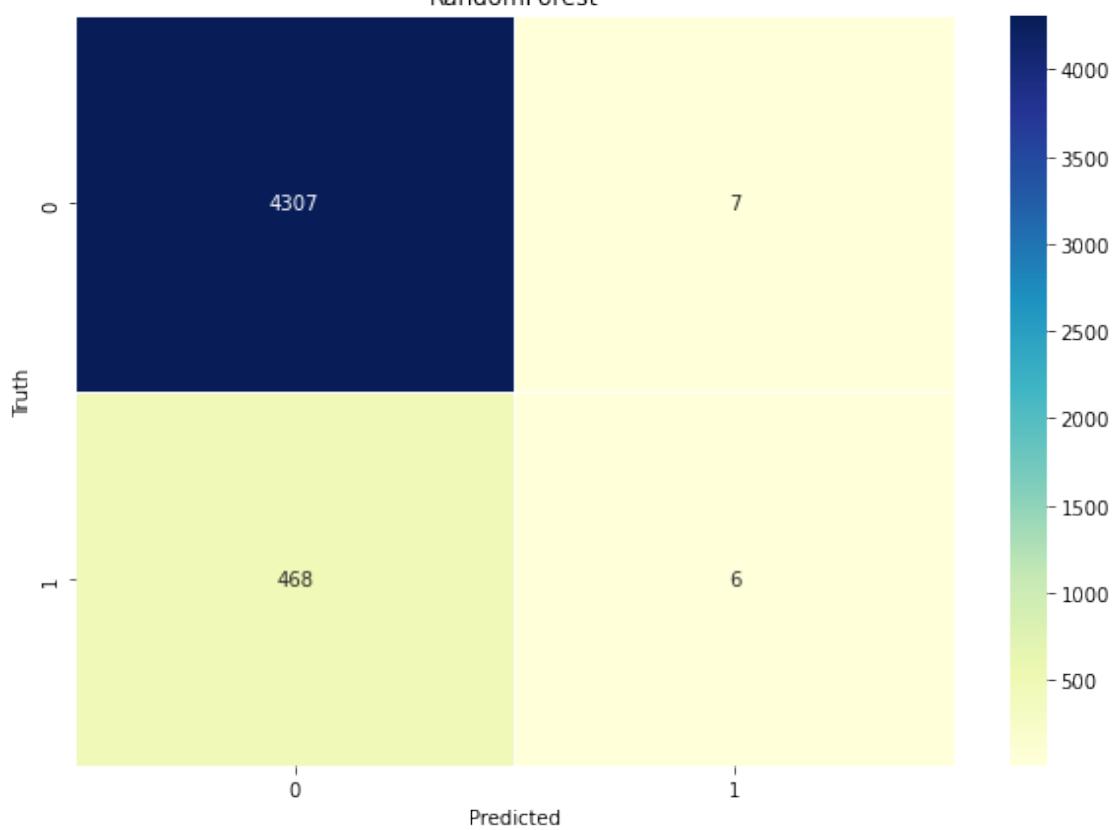
/Users/jatindersinghmalhi/opt/anaconda3/envs/tensorflow/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

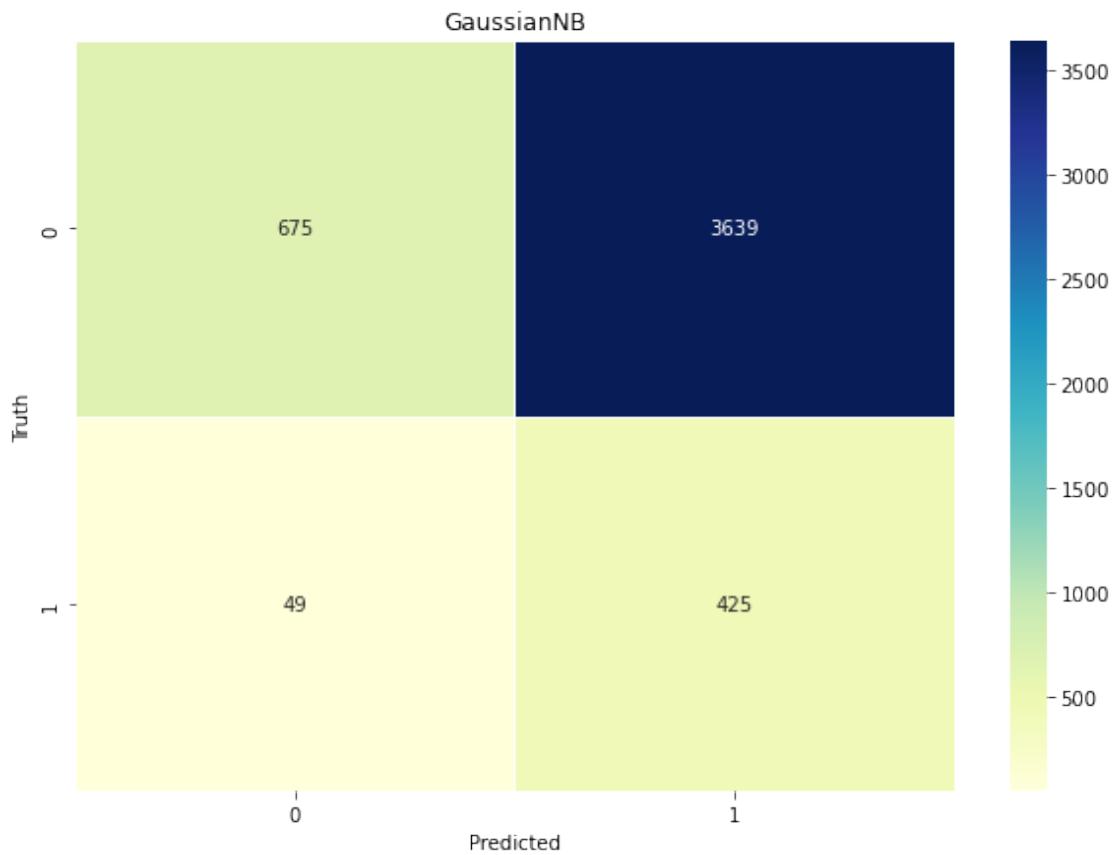
```

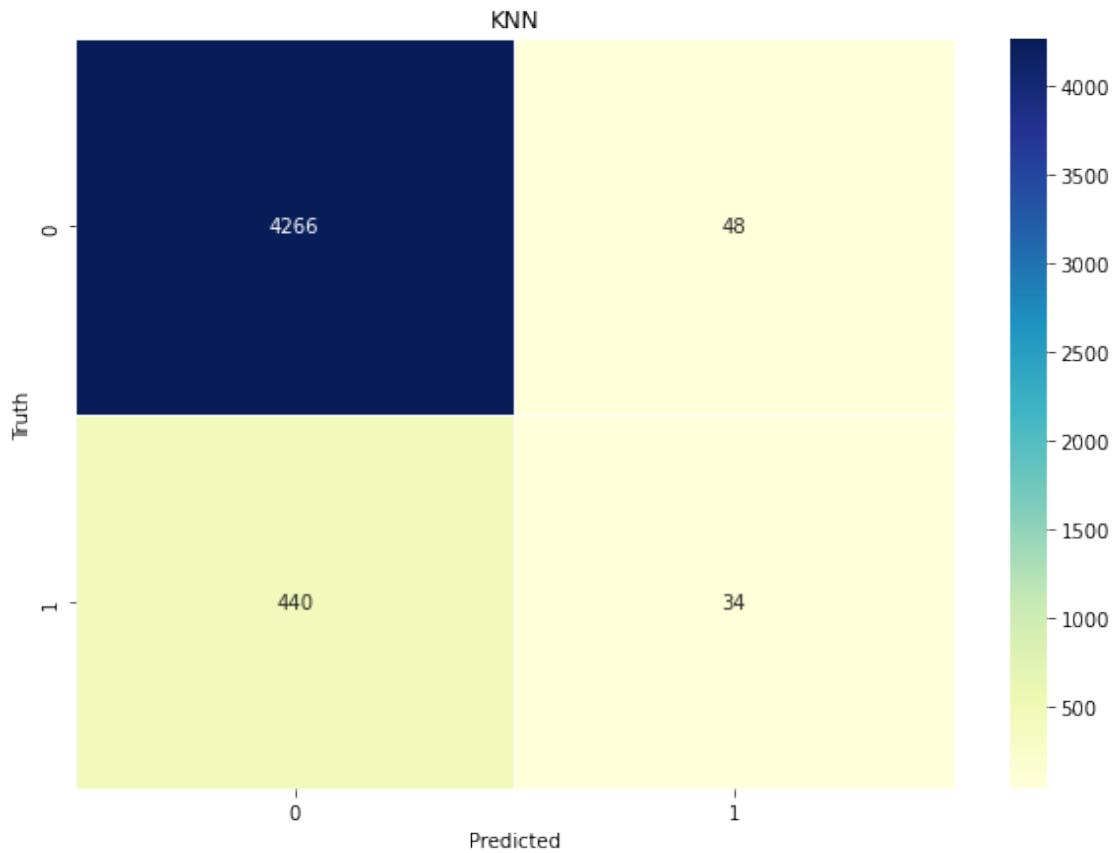




RandomForest







```
[ ]:      model  accuracy  precision  recall
0  LogisticRegression    0.90      0.00     0.00
1      DecisionTree       0.82      0.15     0.18
2      RandomForest       0.90      0.46     0.01
3      GaussianNB        0.23      0.10     0.90
4          KNN            0.90      0.41     0.07
```